

# Interpreting\_your\_Data

September 30, 2018

## 1 Random Forest Model interpretation

```
In [1]: %load_ext autoreload
        %autoreload 2
```

```
In [47]: %matplotlib inline
```

```
from fastai.imports import *
from fastai.structured import *
from pandas_summary import DataFrameSummary
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier
from IPython.display import display
from sklearn import metrics
```

```
In [48]: set_plot_sizes(12,14,16)
```

### 1.1 Load in our data

```
In [49]: # point path to data
        PATH = "data/bulldozers/"
```

```
In [50]: # notice parse_dates, this is an important step when you have a
        # date column thats not already split into its constituent parts
        df_raw = pd.read_csv(f'{PATH}Train.csv', low_memory=False,
                             parse_dates=["saledate"])
```

```
In [51]: df_raw.SalePrice = np.log(df_raw.SalePrice)
        add_datepart(df_raw, 'saledate')
        df_raw.saleYear.head()
```

```
Out[51]: 0    2006
        1    2004
        2    2004
        3    2011
        4    2009
        Name: saleYear, dtype: int64
```

```

In [52]: # train_cats is our solution
# there is some nuance here however
# when assigning categories it is important to be consistent
# notice how this is called "train"_cats
# this is for the training set
# when you try and train your model you typically have a train and valid set
# so 'high' might be 0 in train but it could be 2 in valid
# for future reference this is why fastai has a apply_cats(df, trn)
# you can pass the data frame of the validation set as well as the
# applied categories from the train set to achieve aforementioned consistency
train_cats(df_raw)
#apply_cats(df_raw)

In [53]: df_trn, y_trn, nas = proc_df(df_raw, 'SalePrice')

In [54]: def split_vals(a,n): return a[:n], a[n:]
n_valid = 12000
n_trn = len(df_trn)-n_valid
X_train, X_valid = split_vals(df_trn, n_trn)
y_train, y_valid = split_vals(y_trn, n_trn)
raw_train, raw_valid = split_vals(df_raw, n_trn)

In [55]: def rmse(x,y): return math.sqrt((x-y)**2).mean())

def print_score(m):
    res = [rmse(m.predict(X_train), y_train), rmse(m.predict(X_valid), y_valid),
            m.score(X_train, y_train), m.score(X_valid, y_valid)]
    if hasattr(m, 'oob_score_'): res.append(m.oob_score_)
    print(res)

In [56]: df_raw.head()

```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	\
0	1139246	11.097410	999089	3157	121	3.0	2004	
1	1139248	10.950807	117657	77	121	3.0	1996	
2	1139249	9.210340	434808	7009	121	3.0	2001	
3	1139251	10.558414	1026470	332	121	3.0	2001	
4	1139253	9.305651	1057373	17311	121	3.0	2007	

	MachineHoursCurrentMeter	UsageBand	fiModelDesc	...	saleDay	\
0	68.0	Low	521D	...	16	
1	4640.0	Low	950FII	...	26	
2	2838.0	High	226	...	26	
3	3486.0	High	PC120-6E	...	19	
4	722.0	Medium	S175	...	23	

	saleDayofweek	saleDayofyear	saleIs_month_end	saleIs_month_start	\
0	3	320	False	False	
1	4	86	False	False	

2	3	57	False	False
3	3	139	False	False
4	3	204	False	False

	saleIs_quarter_end	saleIs_quarter_start	saleIs_year_end	saleIs_year_start	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	saleElapsed
0	1163635200
1	1080259200
2	1077753600
3	1305763200
4	1248307200

[5 rows x 65 columns]

## 2 Confidence based on tree variance

For model interpretation, there's no need to use the full dataset on each tree - using a subset will be both faster, and also provide better interpretability (since an overfit model will not provide much variance across trees).

```
In [57]: set_rf_samples(50000)
```

```
In [58]: # the validation set is a little worse due to being all from the future
m = RandomForestRegressor(n_estimators=40, min_samples_leaf=3, max_features=0.5, n_jobs=-1)
m.fit(X_train, y_train)
print_score(m)
```

```
[0.20649529124100546, 0.24944018378023106, 0.910884113505606, 0.8888829227135021, 0.8937400278]
```

We saw how the model averages predictions across the trees to get an estimate - but how can we know the confidence of the estimate? One simple way is to use the standard deviation of predictions, instead of just the mean. This tells us the *relative* confidence of predictions - that is, for rows where the trees give very different results, you would want to be more cautious of using those results, compared to cases where they are more consistent. Using the same example as in the last lesson when we looked at bagging:

```
In [59]: %time preds = np.stack([t.predict(X_valid) for t in m.estimators_])
np.mean(preds[:,0]), np.std(preds[:,0])
```

```
CPU times: user 1.38 s, sys: 40 ms, total: 1.42 s
Wall time: 1.43 s
```

```
Out [59]: (9.250535703082043, 0.2822839234241768)
```

When we use python to loop through trees like this, we're calculating each in series, which is slow! We can use parallel processing to speed things up:

```
In [60]: def get_preds(t): return t.predict(X_valid)
         %time preds = np.stack(parallel_trees(m, get_preds))
         np.mean(preds[:,0]), np.std(preds[:,0])
```


```
CPU times: user 88 ms, sys: 156 ms, total: 244 ms
```

```
Wall time: 511 ms
```

```
Out [60]: (9.250535703082043, 0.2822839234241768)
```

We can see that different trees are giving different estimates this this auction. In order to see how prediction confidence varies, we can add this into our dataset.

```
In [61]: x = raw_valid.copy()
         x['pred_std'] = np.std(preds, axis=0)
         x['pred'] = np.mean(preds, axis=0)
         x.Enclosure.value_counts().plot.barh();
```




```
In [62]: flds = ['Enclosure', 'SalePrice', 'pred', 'pred_std']
         enc_summ = x[flds].groupby('Enclosure', as_index=False).mean()
         enc_summ
```


```
Out [62]:
```

	Enclosure	SalePrice	pred	pred_std
0	EROPS	9.849178	9.840480	0.278144
1	EROPS AC	NaN	NaN	NaN
2	EROPS w AC	10.623971	10.575259	0.274543
3	NO ROPS	NaN	NaN	NaN
4	None or Unspecified	NaN	NaN	NaN
5	OROPS	9.682064	9.684958	0.218898


```
In [63]: enc_summ = enc_summ[~pd.isnull(enc_summ.SalePrice)]
enc_summ.plot('Enclosure', 'SalePrice', 'barh', xlim=(0,11));
```



```
In [64]: enc_summ.plot('Enclosure', 'pred', 'barh', xerr='pred_std', alpha=0.6, xlim=(0,11));
```



```
In [65]: raw_valid.ProductSize.value_counts().plot.barh();
```



```
In [66]: flds = ['ProductSize', 'SalePrice', 'pred', 'pred_std']
         summ = x[flds].groupby(flds[0]).mean()
         summ
```

```
Out[66]:
```

	SalePrice	pred	pred_std
ProductSize			
Compact	9.735093	9.858105	0.344963
Large	10.470589	10.386889	0.351798
Large / Medium	10.691871	10.646844	0.308542
Medium	10.681511	10.614658	0.289467
Mini	9.535147	9.567300	0.252591
Small	10.324448	10.324686	0.313421

```
In [67]: (summ.pred_std/summ.pred).sort_values(ascending=False)
```

```
Out[67]: ProductSize
Compact      0.034993
Large        0.033869
Small        0.030356
Large / Medium 0.028980
Medium       0.027270
Mini         0.026401
dtype: float64
```