

Multi-label Classification

September 7, 2018

0.1 Multi-label Classification (Planet Dataset)

```
In [10]: %reload_ext autoreload
         %autoreload 2
         %matplotlib inline
```

```
In [11]: # import convolution tools
         from fastai.conv_learner import *
         # import pandas for csv visualization
         import pandas as pd
```

```
In [12]: # set our path to the planet data set
         PATH = 'data/planet/'
```

```
In [13]: # list our sub directories
         !ls {PATH}
```

```
models          test-jpg  tmp    train-jpg  train_v2.csv
```

```
In [14]: df = pd.read_csv("data/planet/train_v2.csv")
         print(df.head())
```

	image_name	tags
0	train_0	haze primary
1	train_1	agriculture clear primary water
2	train_2	clear primary
3	train_3	clear primary
4	train_4	agriculture clear habitation primary road

```
In [15]: # plotting library from fastai
         from fastai.plots import *
```

```
In [16]: # gets the first picture in a path
         def get_1st(path):
             # glob lets you return potentially empty file locs without throwing an error
             # this just says give me the first of anything.anything in PATH
             return glob(f'{path}/*.*)')[0]
```

```
In [17]: # go to the valid dogs and cats this is to showcase the difference between single and
# classification
dc_path = "data/dogscats/valid/"

# creates a list of things we will later plot
list_paths = [get_1st(f"{dc_path}cats"), get_1st(f"{dc_path}dogs")]

# plots the list above with sub and main titles
plots_from_files(list_paths, titles=["cat", "dog"], maintitle="Single-label classification")
```

Single-label classification



Remember to use A sigmoid instead of Softmax for multilabel classification


```
In [18]: # this does the same as above
list_paths = [f"{PATH}train-jpg/train_0.jpg", f"{PATH}train-jpg/train_1.jpg"]

titles=["haze primary", "agriculture clear primary water"]

plots_from_files(list_paths, titles=titles, maintitle="Multi-label classification")
```

Multi-label classification

haze primary



agriculture clear primary water



```
In [19]: from fastai.imports import *
         from fastai.transforms import *
         from fastai.dataset import *
         from sklearn.metrics import fbeta_score
         import warnings

         # this a predefined metric by Jeremy for checking error using beta score
         def f2(preds, targs, start=0.17, end=0.24, step=0.01):
             with warnings.catch_warnings():
                 warnings.simplefilter("ignore")
                 return max([fbeta_score(targs, (preds>th), 2, average='samples')
                             for th in np.arange(start,end,step)])

In [20]: # f2 is like a confusion matrix
         # there are lots of different ways you can turn that confusion matrix into a score
         # f2 is f beta with beta=2, the kaggle competition wanted this
         #
         metrics=[f2]
         f_model = resnet34

In [21]: # The csv file containing our labels
         label_csv = f'{PATH}train_v2.csv'

         #o = open(label_csv)
         #li = list(o)
         #n = len(li)-1
```

```
n = len(list(open(label_csv)))-1

val_idx = get_cv_idx(n)
```

```
In [22]: # takes in the size of the image and returns a ImageClassifierData object
def get_data(sz):
    # apply transforms
    tfms = tfms_from_model(f_model, sz, aug_tfms=transforms_top_down, max_zoom=1.05)
    return ImageClassifierData.from_csv(PATH, 'train-jpg', label_csv, tfms=tfms,
                                       suffix='.jpg', val_idx=val_idx, test_name='test-jpg')
```

```
In [23]: # the images are 256 x 256
data = get_data(256)
```

```
In [24]: # ds is data set
# this will give you a single image back

# dl is data loader
# a data loader gives you back a minibatch
# specifically a transformed minibatch
# also you can only get the next minibatch so a good practice
# would be to loop through and grab a minibatch at a time

# for more info google pytorch data set and data loader
```

```
x,y = next(iter(data.val_dl))
```

```
In [25]: # batch size by default is 64 in the function
# so its 64 by 17 possible classes or labels
# this also seems to be transposed so its index x class
y
```

Out[25]:

```


1      0      0 ...      0      0      1
1      0      0 ...      0      0      0
0      0      0 ...      0      0      0
...
0      0      0 ...      0      0      1
1      0      0 ...      0      0      0
1      0      0 ...      0      0      0
[torch.cuda.FloatTensor of size 64x17 (GPU 0)]
```

```
In [26]: # zip will combine lists
# so you will get the 0th thing from the first list and the 0th thing from the second
list(zip(data.classes, y[0]))
```

```
Out[26]: [('agriculture', 1.0),
          ('artisinal_mine', 0.0),
```

```
(
    ('bare_ground', 0.0),
    ('blooming', 0.0),
    ('blow_down', 0.0),
    ('clear', 1.0),
    ('cloudy', 0.0),
    ('conventional_mine', 0.0),
    ('cultivation', 0.0),
    ('habitation', 0.0),
    ('haze', 0.0),
    ('partly_cloudy', 0.0),
    ('primary', 1.0),
    ('road', 0.0),
    ('selective_logging', 0.0),
    ('slash_burn', 0.0),
    ('water', 1.0)]
```

```
In [27]: plt.imshow(data.val_ds.denorm(to_np(x))[0]*1.4);
```



```
In [28]: # we do this here since there is nothing like satellite images in the imagenet model
         # we wouldn't do this with the cats and dogs because we would destroy the weights
         sz=64
```

```
In [29]: # create data object
         data = get_data(sz)
```

```
In [30]: # this speeds things up by saying
         # I won't be needing anything bigger than sz*1.3
```

```

data = data.resize(int(sz*1.3), 'tmp')

HBox(children=(IntProgress(value=0, max=6), HTML(value='')))


In [31]: learn = ConvLearner.pretrained(f_model, data, metrics=metrics)

In [32]: lrf=learn.lr_find()
         learn.sched.plot()

HBox(children=(IntProgress(value=0, description='Epoch', max=1), HTML(value='')))

```

epoch	trn_loss	val_loss	f2
0	0.223709	0.269339	0.825095



```

In [33]: lr = 0.2

In [34]: learn.fit(lr, 3, cycle_len=1, cycle_mult=2)

```

```
HBox(children=(IntProgress(value=0, description='Epoch', max=7), HTML(value='')))
```

epoch	trn_loss	val_loss	f2
0	0.147102	0.133033	0.883338
1	0.141103	0.127185	0.889755
2	0.139051	0.125667	0.892151
3	0.137805	0.124761	0.892294
4	0.136067	0.122831	0.89526
5	0.132859	0.121276	0.896396
6	0.131545	0.121369	0.895219

```
Out [34]: [array([0.12137]), 0.8952189648278935]
```

```
In [35]: # notice the learning rate is greater for all 3 layers compared to the cats and dogs
# this is because imagenet was more suited for the cats and dogs model
# satellite data is much different so we set a higher learning rate
# at lower layers, but still less than our top layers

# Also notice how we don't do this right away
# this is because our layers would be random
# we should train our layers first then unfreeze the imagenet layers
# for retraining
lrs = np.array([lr/9,lr/3,lr])
```

```
In [36]: learn.unfreeze()
learn.fit(lrs, 3, cycle_len=1, cycle_mult=2)
```


```
HBox(children=(IntProgress(value=0, description='Epoch', max=7), HTML(value='')))
```

epoch	trn_loss	val_loss	f2
0	0.120365	0.109936	0.907705
1	0.116826	0.107446	0.907156
2	0.107001	0.101385	0.915106
3	0.112509	0.102195	0.913443
4	0.105939	0.101563	0.915094
5	0.102347	0.097551	0.917749
6	0.097418	0.097372	0.918057

```
Out [36]: [array([0.09737]), 0.9180571787743648]
```

```
In [37]: learn.save(f'{sz}')
```

```
In [38]: learn.sched.plot_loss()
```



```
In [39]: sz=128
```

```
In [40]: learn.set_data(get_data(sz))
         learn.freeze()
         learn.fit(lr, 3, cycle_len=1, cycle_mult=2)
```

```
HBox(children=(IntProgress(value=0, description='Epoch', max=7), HTML(value='')))
```

epoch	trn_loss	val_loss	f2
0	0.09867	0.094844	0.91859
1	0.097065	0.093558	0.920441
2	0.095911	0.092839	0.920789
3	0.09785	0.092624	0.921624
4	0.098201	0.092687	0.920607
5	0.09571	0.091823	0.92184
6	0.094101	0.09177	0.921351

```
Out[40]: [array([0.09177]), 0.9213513015617453]
```

```
In [41]: learn.unfreeze()
         learn.fit(lrs, 3, cycle_len=1, cycle_mult=2)
         learn.save(f'{sz}')
```



```
HBox(children=(IntProgress(value=0, description='Epoch', max=7), HTML(value='')))
```

epoch	trn_loss	val_loss	f2
0	0.094061	0.087197	0.926153
1	0.092818	0.086996	0.928174
2	0.091056	0.084689	0.929947
3	0.091446	0.089575	0.92483
4	0.087511	0.085343	0.929945
5	0.085561	0.08384	0.930323
6	0.083966	0.083565	0.930993

```
In [47]: sz=256
```

```
In [43]: learn.set_data(get_data(sz))
         learn.freeze()
         learn.fit(lr, 3, cycle_len=1, cycle_mult=2)
```

```
HBox(children=(IntProgress(value=0, description='Epoch', max=7), HTML(value='')))
```

epoch	trn_loss	val_loss	f2
0	0.093788	0.088406	0.925103
1	0.092152	0.087738	0.926359
2	0.089541	0.086787	0.927078
3	0.090457	0.087079	0.925858
4	0.088064	0.086749	0.926517
5	0.091159	0.085939	0.927563
6	0.08619	0.085856	0.92778

```
Out[43]: [array([0.08586]), 0.9277797506975684]
```

```
In [44]: learn.unfreeze()
         learn.fit(lrs, 3, cycle_len=1, cycle_mult=2)
         learn.save(f'{sz}')

```

```
HBox(children=(IntProgress(value=0, description='Epoch', max=7), HTML(value='')))
```

epoch	trn_loss	val_loss	f2
0	0.087228	0.082125	0.931622
1	0.090669	0.083097	0.929711
2	0.083937	0.08099	0.932874
3	0.087432	0.083413	0.930149
4	0.083918	0.082244	0.932266

5	0.079948	0.08139	0.932267
6	0.077498	0.080706	0.933599

```
In [45]: multi_preds, y = learn.TTA()  
        preds = np.mean(multi_preds, 0)
```

```
In [46]: f2(preds,y)
```

```
Out[46]: 0.9316254029148617
```