

SloanDigitalSkySurvey

September 30, 2018

1 Sloan Digital Sky Survey

```
In [79]: from fastai.imports import *
        from fastai.structured import *

        from pandas_summary import DataFrameSummary
        from sklearn.ensemble import RandomForestClassifier
        from IPython.display import display

        from sklearn import metrics
```

```
In [80]: # set path to data
        PATH = "/home/chris/Datasets/SkySurvey/"
```

```
In [81]: # take a look to make sure it's there
        !ls {PATH}
```

```
'Skyserver_SQL2_27_2018 6_51_39 PM.csv.zip'    Train.csv
```

```
In [82]: # create a pandas df for our data
        df_raw = pd.read_csv(f'{PATH}Train.csv', low_memory=False)
```

```
In [83]: df_raw.head()
```

```
Out[83]:
```

	objid	ra	dec	u	g	r	i	\
0	1.237650e+18	183.531326	0.089693	19.47406	17.04240	15.94699	15.50342	
1	1.237650e+18	183.598371	0.135285	18.66280	17.21449	16.67637	16.48922	
2	1.237650e+18	183.680207	0.126185	19.38298	18.19169	17.47428	17.08732	
3	1.237650e+18	183.870529	0.049911	17.76536	16.60272	16.16116	15.98233	
4	1.237650e+18	183.883288	0.102557	17.55025	16.26342	16.43869	16.55492	

	z	run	rerun	camcol	field	specobjid	class	redshift	plate	\
0	15.22531	752	301	4	267	3.722360e+18	STAR	-0.000009	3306	
1	16.39150	752	301	4	267	3.638140e+17	STAR	-0.000055	323	
2	16.80125	752	301	4	268	3.232740e+17	GALAXY	0.123111	287	
3	15.90438	752	301	4	269	3.722370e+18	STAR	-0.000111	3306	
4	16.61326	752	301	4	269	3.722370e+18	STAR	0.000590	3306	

	mjd	fiberid
0	54922	491
1	51615	541
2	52023	513
3	54922	510
4	54922	512

In [84]: `df_raw.dtypes`

```
Out[84]: objid      float64
         ra         float64
         dec        float64
         u          float64
         g          float64
         r          float64
         i          float64
         z          float64
         run        int64
         rerun      int64
         camcol     int64
         field      int64
         specobjid  float64
         class      object
         redshift   float64
         plate      int64
         mjd        int64
         fiberid    int64
         dtype: object
```

In [85]: *# our data has 10,000 rows and 18 columns*
`df_raw.shape`

```
Out[85]: (10000, 18)
```

In [86]: *#df_raw = df_raw.drop('objid', axis=1)*
`df = train_cats(df_raw)`

In [87]: *# GALAXY=0 QSO=1 STAR=2*
`df_raw['class'].cat.categories`

```
Out[87]: Index(['GALAXY', 'QSO', 'STAR'], dtype='object')
```

In [88]: *# I chose to show 20 so we can see every type of object*
`df_raw['class'].cat.codes.head(20)`

```
Out[88]: 0      2
         1      2
         2      0
```

```

3      2
4      2
5      2
6      0
7      2
8      2
9      0
10     2
11     2
12     2
13     2
14     0
15     1
16     2
17     1
18     2
19     0
dtype: int8

```

```
In [89]: # split our independent and dependent variables
```

```

X_train = df_raw.drop(['class'], axis=1)
y_train = df_raw['class']

```

```
# instantiate our classifier
```

```
m = RandomForestClassifier(n_jobs=-1)
```

```
# fit our data
```

```
m.fit(X_train, y_train)
```

```

Out[89]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=-1,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)

```

```
In [90]: # Let's see our training accuracy
```

```
m.score(X_train, y_train)
```

```
Out[90]: 0.9983
```

```
In [91]: # possible classes
```

```
m.classes_
```

```
Out[91]: array(['GALAXY', 'QSO', 'STAR'], dtype=object)
```

```
In [92]: # print this out to compare with processed df
```

```
y_train[:5]
```

```
Out [92]: 0      STAR
          1      STAR
          2    GALAXY
          3      STAR
          4      STAR
          Name: class, dtype: category
          Categories (3, object): [GALAXY < QSO < STAR]
```

```
In [93]: # process dataframe
          df, y, nas = proc_df(df_raw, 'class')
```

```
In [94]: # our vectorized dependent variable
          y[:5]
```

```
Out [94]: array([2, 2, 0, 2, 2], dtype=int8)
```

```
In [95]: df_raw.head()
```

```
Out [95]:
```

	objid	ra	dec	u	g	r	i	\
0	1.237650e+18	183.531326	0.089693	19.47406	17.04240	15.94699	15.50342	
1	1.237650e+18	183.598371	0.135285	18.66280	17.21449	16.67637	16.48922	
2	1.237650e+18	183.680207	0.126185	19.38298	18.19169	17.47428	17.08732	
3	1.237650e+18	183.870529	0.049911	17.76536	16.60272	16.16116	15.98233	
4	1.237650e+18	183.883288	0.102557	17.55025	16.26342	16.43869	16.55492	

	z	run	rerun	camcol	field	specobjid	class	redshift	plate	\
0	15.22531	752	301	4	267	3.722360e+18	STAR	-0.000009	3306	
1	16.39150	752	301	4	267	3.638140e+17	STAR	-0.000055	323	
2	16.80125	752	301	4	268	3.232740e+17	GALAXY	0.123111	287	
3	15.90438	752	301	4	269	3.722370e+18	STAR	-0.000111	3306	
4	16.61326	752	301	4	269	3.722370e+18	STAR	0.000590	3306	

	mjd	fiberid
0	54922	491
1	51615	541
2	52023	513
3	54922	510
4	54922	512

```
In [96]: # this function will split the dataframe 80% train 20% valid
          def split_vals(a,n):
              return a[:n].copy(), a[n:].copy()

          n_valid = 2000
          n_train = len(df_raw)-n_valid

          # split the data initially
          raw_train, raw_valid = split_vals(df_raw, n_train)
```

```

    # now split each of those
    X_train, X_valid = split_vals(df, n_train)
    y_train, y_valid = split_vals(y, n_train)

    X_train.shape, y_train.shape, X_valid.shape, y_valid.shape

Out[96]: ((8000, 17), (8000,), (2000, 17), (2000,))

In [97]: # y_train took on traits of y
         y_train

Out[97]: array([2, 2, 0, ..., 0, 1, 0], dtype=int8)

In [98]: def print_score(m):
         res = [m.score(X_train, y_train), m.score(X_train, y_train)]
         if hasattr(m, 'oob_score_'):
             res.append(m.oob_score_)
         print(res)

In [99]: m = RandomForestClassifier(n_jobs=-1)
         %time m.fit(X_train, y_train)
         print_score(m)

CPU times: user 294 ms, sys: 8.05 ms, total: 302 ms
Wall time: 243 ms
[0.99825, 0.99825]

In [100]: # possible classes after changing array
          m.classes_

Out[100]: array([0, 1, 2], dtype=int8)

In [101]: m = RandomForestClassifier(n_estimators=40, n_jobs=-1, oob_score=True)
          m.fit(X_train, y_train)
          print_score(m)

[0.99975, 0.99975, 0.99]

In [102]: m.predict([[1.237650e+18, 130.1993148, 51.00089591, 18.68415, 17.78766, 17.51412, 17.41314

Out[102]: array([2], dtype=int8)

In [103]: m.classes_

Out[103]: array([0, 1, 2], dtype=int8)

In [104]: m.feature_importances_

```

```
Out[104]: array([0.          , 0.00543, 0.0056 , 0.01543, 0.04344, 0.03159, 0.06321, 0.04574, 0.003
0.00374, 0.11626, 0.47978, 0.08179, 0.09869, 0.00452])
```

```
In [105]: X_train.head()
```

```
Out[105]:
```

	objid	ra	dec	u	g	r	i	\
0	1.237650e+18	183.531326	0.089693	19.47406	17.04240	15.94699	15.50342	
1	1.237650e+18	183.598371	0.135285	18.66280	17.21449	16.67637	16.48922	
2	1.237650e+18	183.680207	0.126185	19.38298	18.19169	17.47428	17.08732	
3	1.237650e+18	183.870529	0.049911	17.76536	16.60272	16.16116	15.98233	
4	1.237650e+18	183.883288	0.102557	17.55025	16.26342	16.43869	16.55492	

	z	run	rerun	camcol	field	specobjid	redshift	plate	mjd	\
0	15.22531	752	301	4	267	3.722360e+18	-0.000009	3306	54922	
1	16.39150	752	301	4	267	3.638140e+17	-0.000055	323	51615	
2	16.80125	752	301	4	268	3.232740e+17	0.123111	287	52023	
3	15.90438	752	301	4	269	3.722370e+18	-0.000111	3306	54922	
4	16.61326	752	301	4	269	3.722370e+18	0.000590	3306	54922	

	fiberid
0	491
1	541
2	513
3	510
4	512

```
In [106]: X_train = X_train.drop('rerun', axis=1)
X_valid = X_valid.drop('rerun', axis=1)
m = RandomForestClassifier(n_estimators=40, n_jobs=-1, oob_score=True)
m.fit(X_train, y_train)
```

```
Out[106]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=40, n_jobs=-1,
oob_score=True, random_state=None, verbose=0, warm_start=False)
```

```
In [107]: print_score(m)
```

```
[0.99975, 0.99975, 0.987875]
```

```
In [108]: X_train = X_train.drop('objid', axis=1)
X_vaild = X_valid.drop('objid', axis=1)

m = RandomForestClassifier(n_estimators=40, n_jobs=-1, oob_score=True)
m.fit(X_train, y_train)
```

```
Out[108]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=40, n_jobs=-1,
                                oob_score=True, random_state=None, verbose=0, warm_start=False)
```

```
In [109]: print_score(m)
```

```
[0.999875, 0.999875, 0.988375]
```

```
In [110]: m.feature_importances_
```

```
Out[110]: array([0.00658, 0.00705, 0.02235, 0.04421, 0.04566, 0.05819, 0.05578, 0.00469, 0.00118,
                 0.42183, 0.10682, 0.12361, 0.00694])
```

```
In [111]: X_train.head()
```

```
Out[111]:
```

	ra	dec	u	g	r	i	z	\
0	183.531326	0.089693	19.47406	17.04240	15.94699	15.50342	15.22531	
1	183.598371	0.135285	18.66280	17.21449	16.67637	16.48922	16.39150	
2	183.680207	0.126185	19.38298	18.19169	17.47428	17.08732	16.80125	
3	183.870529	0.049911	17.76536	16.60272	16.16116	15.98233	15.90438	
4	183.883288	0.102557	17.55025	16.26342	16.43869	16.55492	16.61326	

	run	camcol	field	specobjid	redshift	plate	mjd	fiberid
0	752	4	267	3.722360e+18	-0.000009	3306	54922	491
1	752	4	267	3.638140e+17	-0.000055	323	51615	541
2	752	4	268	3.232740e+17	0.123111	287	52023	513
3	752	4	269	3.722370e+18	-0.000111	3306	54922	510
4	752	4	269	3.722370e+18	0.000590	3306	54922	512

```
In [112]: m.oob_score_
```

```
Out[112]: 0.988375
```